Research Statement — Dani Yogatama

The ability to continuously learn and generalize to new problems quickly is a hallmark of general intelligence. In my research, I design **machine learning** models to advance artificial intelligence on this front. Within machine learning, my primary application area is **natural language processing**. While deep learning has driven progress in many language understanding tasks, existing models have been shown to require a lot of in-domain training examples, rapidly overfit to particular datasets, and are prone to catastrophic forgetting (Yogatama *et al.*, 2019). In order to make progress toward artificial agents that exhibit general linguistic intelligence, we need machine learning models that can (i) deal with the full complexity of natural language across a variety of tasks, (ii) effectively store and reuse representations, combinatorial modules, and previously acquired linguistic knowledge, and (iii) adapt to new tasks in new environments with little experience. My research seeks to answer the following questions:

- How do we find the best way to represent language for various tasks? (§1)

- How do we store and reuse linguistic knowledge to avoid catastrophic forgetting? (§2)

- How do we ensure sample efficient generalization to new problems? (§3)

I take inspirations from cognitive and neuroscience and use techniques from deep learning, probabilistic graphical models, information theory, and many others to answer these questions. I have broad interests in many aspects of machine learning and natural language processing. In §4, I discuss future directions.

1 Representation Learning

The performance of a machine learning model heavily depends on how the data is represented in the model. For example, when working with text data, we can represent it as a sequence of words, subwords, or characters. Furthermore, each textual unit can be represented as strings, binary vectors, or real vectors. Traditional methods rely on human expertise to choose the best representation. In my research, I contribute theoretical foundations and analytical insights and use them to improve representation learning methods.

In Kong *et al.* (2019a), we showed that state-of-the-art language representation learning methods maximize an objective function that is a lower bound on the mutual information between different parts of a word sequence. Our formulation provides an information theoretic perspective that unifies classical word embedding models and modern contextual embeddings. It also leads to a principled framework that can be used to construct new self-supervised tasks. The resulting framework offers a holistic view of representation learning methods to transfer knowledge and translate progress across multiple domains (e.g., natural language processing, computer vision, audio processing).

Human language learning is facilitated by abstractions that are independent of any particular language. In Artetxe *et al.* (2019), we presented a method to transfer a representation learning model for a particular language (e.g., English) to other languages. We evaluated the model in a zero-shot setting (without labeled training data in the new languages) and demonstrated that a language model trained on English learns generalizable abstractions that are reusable in other languages. As a part of this project, we also created a new multilingual question answering dataset that was released to the research community.

In order to build effective language representations, we need combinatorial modules which compose words into representations of phrases, sentences, and documents. In Yogatama *et al.* (2017b) and Maillard *et al.* (2019), we explored two methods based on reinforcement learning and differentiable parsers that compute representations of the meaning

of sentences by composing representations of words and phrases. We showed that our automatic approaches yield to better representations compared to methods that rely on explicit supervisions (e.g., syntactic parse trees of sentences).

My other work in this area includes variational language models (Kong *et al.*, 2019b), learning ordered and sparse word representations (Yogatama *et al.*, 2015b; Faruqui *et al.*, 2015) and using Bayesian optimization for representation learning (Yogatama *et al.*, 2015a).

2   Memory Models

Short-term and long-term memory systems is an integral part of human intelligence. I work on combining neural networks with memory modules to make artificial agents that can store knowledge and reuse it effectively.

State-of-the-art machine learning models work well on a single dataset given enough training examples, but they often fail to isolate and reuse previously acquired knowledge when the data distribution shifts (e.g., when presented with a new dataset)—a phenomenon known as catastrophic forgetting (McCloskey and Cohen, 1989; Ratcliff, 1990). In de Masson d'Autume *et al.* (2019), we presented a method that augments a neural network language model with an episodic memory module to address this problem. We considered a lifelong learning setup where the model continuously learns from a stream of examples from multiple datasets. The memory module is used to store previously seen examples, which are then used for sparse experience replay and local adaptation (such a process bears some similarity to memory consolidation in human learning; McGaugh 2000). We showed that the episodic memory module mitigates catastrophic forgetting and enables the model to accumulate knowledge throughout its lifetime in question answering and text classification experiments.

In Yogatama *et al.* (2018), we compared and analyzed several working memory architectures (i.e., sequential, random access, and stack memory architectures) that are used to capture linguistic dependencies when learning a language model. We also proposed a generalization to existing continuous stack memory models. We observed that stack-based architectures that encode a bias resembling hierarchical dependencies inherently found in natural language perform the best in explaining the distribution of words in natural language. This result is in line with linguistic theories that claim a context-free backbone for natural language (Chomsky, 1957).

3   Sample Efficient Learning

Neural networks perform well at pattern recognition, but they require a large number of in-domain training examples. I strive to build agents that exhibit sample efficient generalization.

Existing models are typically evaluated by their performances on a held-out test set for a task of interest. The metrics used for evaluation capture an essential aspect of intelligence: being able to generalize from experience with a class of inputs to new inputs. However, they are incomplete as none of them asses a defining attribute of general intelligence: the ability to generalize rapidly to a new task. In Yogatama *et al.* (2019), we proposed a new metric—online codelength—that quantifies how quickly an agent (model) learns a new task. In addition to overall performance, it also rewards models that perform well with limited numbers of training examples. Online codelength is rooted in information theory and is based on connections between generalization, compression, and comprehension (Wolff, 1982; Chaitin, 2007). Our metric can be used across a number of tasks by any *probabilistic* model, allows seamless incorporation of other model and training properties (e.g., model complexity, training cost) if desired, and reflects improvements in generalization performance due to better architecture, better initialization, and better learning algorithms. We also showed that it correlates well with standard evaluation metrics such as accuracy and F_1 scores.

A key reason that existing models are data hungry and generalize poorly to new tasks is that they rely on task specific components that are trained discriminatively, among others. In Yogatama *et al.* (2017a), we empirically characterized the performance of discriminative and generative recurrent neural networks for text classification. We found that generative models that have no task-specific modules approach their asymptotic error rate more rapidly than their discriminative counterparts—the same pattern that Ng and Jordan (2001) proved holds for linear classification models that make more naïve conditional independence assumptions. Building on this finding, we showed that RNN-based generative classification models are more robust to shifts in the data distribution in a series of experiments in zero-shot and continual learning settings.

My earlier work in this area includes a sample efficient hyperparameter tuning method based on transfer learning (Yogatama and Mann, 2014).

4 Future Directions

Achieving general linguistic intelligence requires advances in many areas. I intend to continue working toward this goal and I am especially interested in exploring the following directions.

Semi-parameteric models.   I think that limitations of existing approaches (e.g., data hungry, catastrophic forgetting, overfitting to a dataset instead of solving a task) are inherently caused by the way we train our agents as big parametric models. In my research, I have been exploring methods to combine non-parametric components such as memory modules with parametric neural networks. I am excited about this research direction, both in terms of how to combine these two modules effectively and how to consolidate (compress) past experiences (i.e., learning what to remember and forget) to manage the time and space complexity when using non-parametric components. I plan to continue taking inspirations from neuroscience and cognitive science to make progress in this area.

Hierarchical generative models.   In addition to being more sample efficient, a perfect generative language model—in theory—should be able to do any linguistic task (e.g., by formulating the task as a question and querying the language model to generate answers). Generative modeling is also crucial for imagination and planning. I believe that investing in generative language models would lead to major advances. I am interested in designing hierarchical language models which meta learn from both task and example distributions. For example, we can create a hierarchical model where each example for a task is drawn from a distribution that depends on a task variable (e.g., a task embedding which is a function of examples in the task). Such a model would be robust to data distribution shifts and generalize to new tasks more efficiently.

Unsupervised representation learning.   The ability to learn without explicit supervision is integral to human intelligence. Over the last few years, progress in this area has driven advances in many downstream tasks. The rapid pace of empirical progress created a gap between our theoretical understanding of state-of-the-art models and their practical applications. I think understanding these models is crucial to assess their limitations and provide a launchpad for future breakthroughs. I am eager to continue working on unsupervised and self-supervised learning. In language, a particular weakness that I plan to address is on how existing models fail to learn to encode persistent knowledge that is useful across sentences.

Summary.   I think we should be moving toward a universal lifelong model that is capable of performing multiple tasks and use its experience to continually improve over time. I have been fortunate to work with leaders in this area in my research career, and I am excited to continue working toward my long-term goal to build a general linguistic agent that is capable of understanding and generating natural language.

References

Artetxe, M., Ruder, S., and Yogatama, D. (2019). On the cross-lingual transferability of monolingual representations. *arXiv preprint*.

Chaitin, G. J. (2007). On the intelligibility of the universe and the notions of simplicity, complexity and irreducibility. In *Thinking about Godel and Turing: Essays on Complexity, 1970–2007*. World Scientific.

Chomsky, N. (1957). *Syntactic Structures*. Mouton.

de Masson d'Autume, C., Ruder, S., Kong, L., and Yogatama, D. (2019). Episodic memory in lifelong language learning. In *Proc. of NeurIPS*.

Faruqui, M., Tsvetkov, Y., Yogatama, D., Dyer, C., and Smith, N. A. (2015). Sparse binary word vector representations. In *Proc. of ACL*.

Kong, L., de Masson d'Autume, C., Yu, L., Ling, W., Dai, Z., and Yogatama, D. (2019a). A mutual information maximization perspective of language representation learning. *arXiv preprint*.

Kong, L., Melis, G., Ling, W., Yu, L., and Yogatama, D. (2019b). Variational smoothing in recurrent neural network language models. In *Proc. of ICLR*.

Maillard, J., Clark, S., and Yogatama, D. (2019). Jointly learning sentence embeddings and syntax with unsupervised tree-LSTMs. *Journal of Natural Language Engineering*, **25**(4), 433–449.

McCloskey, M. and Cohen, N. J. (1989). Catastrophic interference in connectionist networks: The sequential learning problem. In *Psychology of Learning and Motivation*, volume 24, pages 109–165. Elsevier.

McGaugh, J. L. (2000). Memory–a century of consolidation. *Science*, **287**(5451), 248–251.

Ng, A. Y. and Jordan, M. I. (2001). On generative and discriminative classifiers: A comparison of logistic regression and naive bayes. In *Proc. of NIPS*.

Ratcliff, R. (1990). Connectionist models of recognition memory: constraints imposed by learning and forgetting functions. *Psychological Review*, **97**(2), 285.

Wolff, J. G. (1982). Language acquisition, data compression and generalization. *Language & Communication*, **2**(1), 57–89.

Yogatama, D. and Mann, G. (2014). Efficient transfer learning method for automatic hyperparameter tuning. In *Proc. of AISTATS*.

Yogatama, D., Kong, L., and Smith, N. A. (2015a). Bayesian optimization of text representations. In *Proc. of EMNLP*.

Yogatama, D., Faruqui, M., Dyer, C., and Smith, N. A. (2015b). Learning word representations with hierarchical sparse coding. In *Proc. of ICML*.

Yogatama, D., Dyer, C., Ling, W., and Blunsom, P. (2017a). Generative and discriminative recurrent neural networks. *arXiv:1703.01898*.

Yogatama, D., Blunsom, P., Dyer, C., Grefenstette, E., and Ling, W. (2017b). Learning to compose words into sentences with reinforcement learning. In *Proc. of ICLR*.

Yogatama, D., Miao, Y., Melis, G., Ling, W., Kuncoro, A., Dyer, C., and Blunsom, P. (2018). Memory architectures in recurrent neural network language models. In *Proc. of ICLR*.

Yogatama, D., de Masson d'Autume, C., Connor, J., Kocisky, T., Chrzanowski, M., Kong, L., Lazaridou, A., Ling, W., Yu, L., Dyer, C., and Blunsom, P. (2019). Learning and evaluating general linguistic intelligence. *arXiv:1901.11373*.